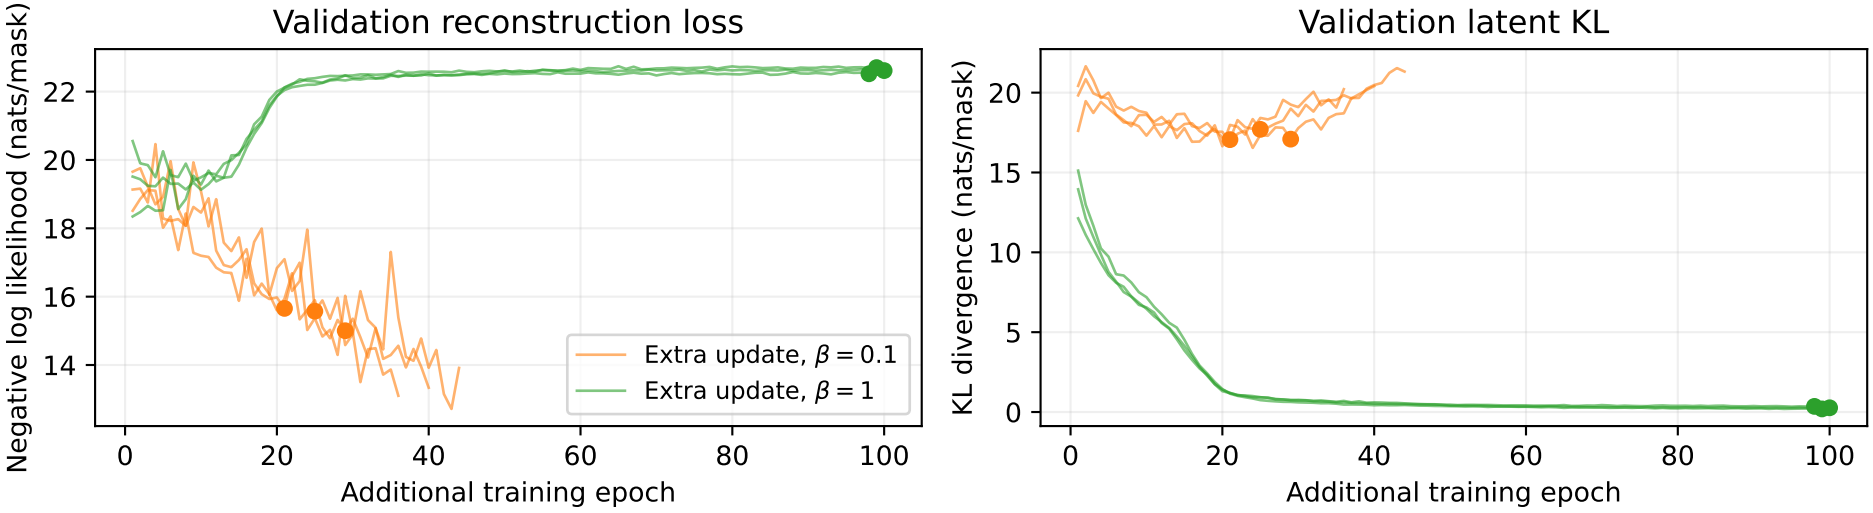
Two adaptive variants on the SAME expanded validation set
One line per fitted seed; dots are the selected checkpoints
Validation reconstruction loss
Negative log likelihood (nats/mask)
Additional training epoch
Extra update, $\beta = 0.1$
Extra update, $\beta = 1$
Validation latent KL
KL divergence (nats/mask)
Additional training epoch